

EDA DIGITAL ASSIGNMENT

BISWAYAN MANDAL - 21BDS0024

Performing Exploration on given auto.csv dataset

1. Loading the Dataset

We use the pandas library to read the dataset from a CSV file into a DataFrame. This allows us to explore and manipulate the data efficiently. We also check the dataset's size (number of rows and columns), column names, data types, a sample of the first few rows, and the count of missing values.

2. Checking for Duplicates

After loading the data, we check for duplicate rows using the `duplicated()` method. Duplicate rows can distort analysis and models, so if any are found, they are removed using `drop_duplicates()`.

3. Summary Statistics

We use the `describe()` method to get key statistical metrics (mean, standard deviation, min, max, and percentiles) for all numerical columns. This gives an initial understanding of the data distribution and highlights potential anomalies like very large or small values.

4. Univariate Analysis

Univariate analysis focuses on one feature at a time:

- We use histograms to visualize the distribution of each numerical column.
- These histograms help identify the central tendency (mean or median), spread (variance or standard deviation), and shape (e.g., skewness) of the data.
- For example, the mpg histogram shows how car fuel efficiencies are distributed.

5. Bivariate Analysis

Bivariate analysis examines relationships between two variables:

- Scatterplots are used to visualize how one numerical feature (e.g., weight) relates to another (e.g., mpg).
- This helps uncover trends like positive, negative, or no correlation.
- For example, weight and mpg might show a negative correlation (heavier cars typically have lower mileage).

6. Multivariate Analysis

Multivariate analysis examines relationships among three or more variables:

- A **correlation heatmap** is used to see how strongly numerical variables are related to one another.
- Before calculating correlations, non-numeric columns (e.g., name) are excluded because they cannot be correlated numerically.
- The heatmap visually represents correlations using color intensity. Strong positive correlations are marked with dark shades (close to 1), while strong negative correlations (close to -1) are marked with the opposite.

7. Handling Errors (ValueError Issue)

The correlation calculation failed initially because the dataset included non-numeric columns like name, which cannot be processed for numerical correlations. To fix this:

- We selected only numeric columns using `select_dtypes()` before calculating correlations.
- This ensures that the correlation matrix contains valid numerical data only.

```
In [1]: import pandas as pd

# Load the dataset
file_path = 'C:/smth/Auto.csv' # Replace with the correct path
data = pd.read_csv(file_path)

# Display basic information
print("Shape of the dataset:", data.shape)
print("\nColumns in the dataset:\n", data.columns)
print("\nData types:\n", data.dtypes)
print("\nFirst few rows of the dataset:\n", data.head())
print("\nMissing values:\n", data.isnull().sum())
```

Shape of the dataset: (392, 9)

Columns in the dataset:

```
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
      'acceleration', 'year', 'origin', 'name'],
      dtype='object')
```

Data types:

```
mpg          float64
cylinders    int64
displacement float64
horsepower   int64
weight       int64
acceleration float64
year         int64
origin       int64
name        object
dtype: object
```

First few rows of the dataset:

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70	
1	15.0	8	350.0	165	3693	11.5	70	
2	18.0	8	318.0	150	3436	11.0	70	
3	16.0	8	304.0	150	3433	12.0	70	
4	17.0	8	302.0	140	3449	10.5	70	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

Missing values:

```
mpg          0
cylinders    0
displacement 0
horsepower   0
weight       0
acceleration 0
year         0
origin       0
name        0
dtype: int64
```

```
In [2]: # Summary statistics
print("\nSummary Statistics:\n", data.describe())

# Check and handle duplicates
duplicates = data.duplicated().sum()
print(f"\nNumber of duplicate rows: {duplicates}")
if duplicates > 0:
    data = data.drop_duplicates()
```

Summary Statistics:

	mpg	cylinders	displacement	horsepower	weight \
count	392.000000	392.000000	392.000000	392.000000	392.000000
mean	23.445918	5.471939	194.411990	104.469388	2977.584184
std	7.805007	1.705783	104.644004	38.491160	849.402560
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.000000	4.000000	105.000000	75.000000	2225.250000
50%	22.750000	4.000000	151.000000	93.500000	2803.500000
75%	29.000000	8.000000	275.750000	126.000000	3614.750000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	year	origin
count	392.000000	392.000000	392.000000
mean	15.541327	75.979592	1.576531
std	2.758864	3.683737	0.805518
min	8.000000	70.000000	1.000000
25%	13.775000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.025000	79.000000	2.000000
max	24.800000	82.000000	3.000000

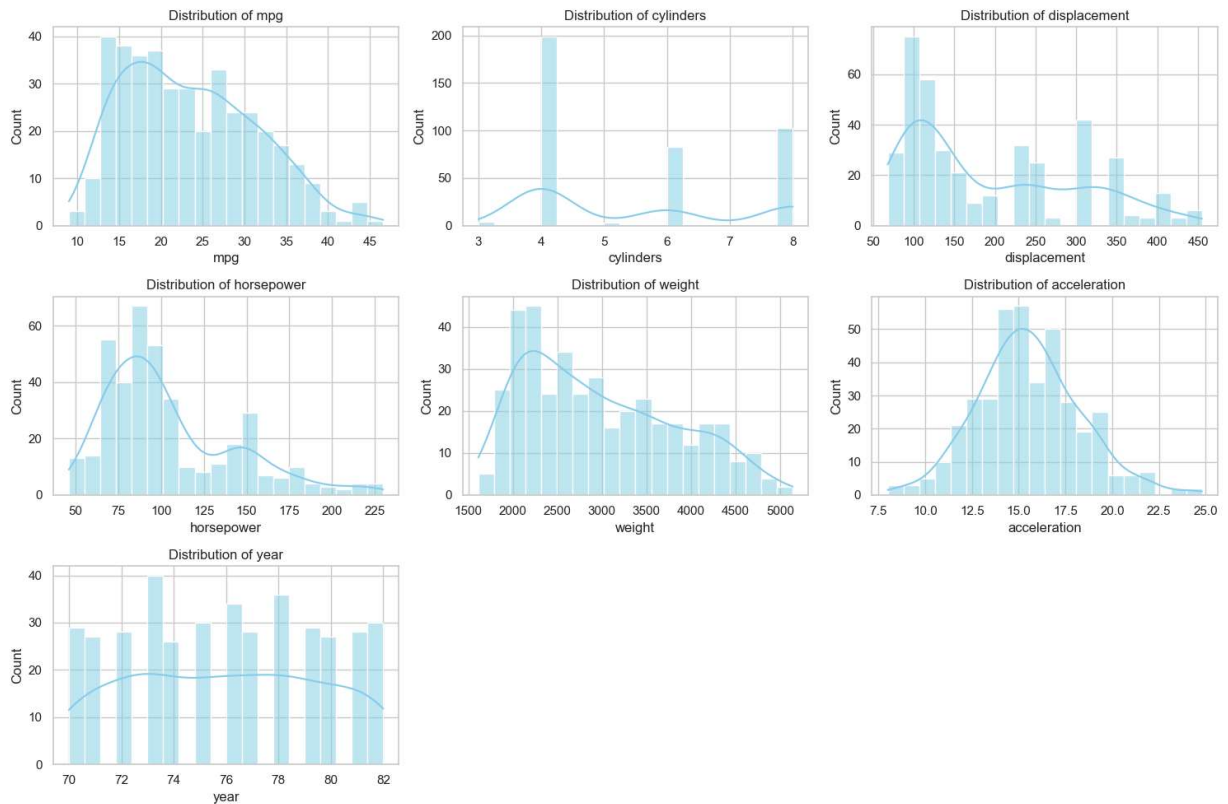
Number of duplicate rows: 0

```
In [3]: import matplotlib.pyplot as plt
import seaborn as sns

sns.set(style="whitegrid")
plt.figure(figsize=(15, 10))

# Plot distributions for all numerical features
numerical_columns = ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'a
for i, column in enumerate(numerical_columns):
    plt.subplot(3, 3, i + 1)
    sns.histplot(data[column], kde=True, bins=20, color='skyblue')
    plt.title(f'Distribution of {column}')

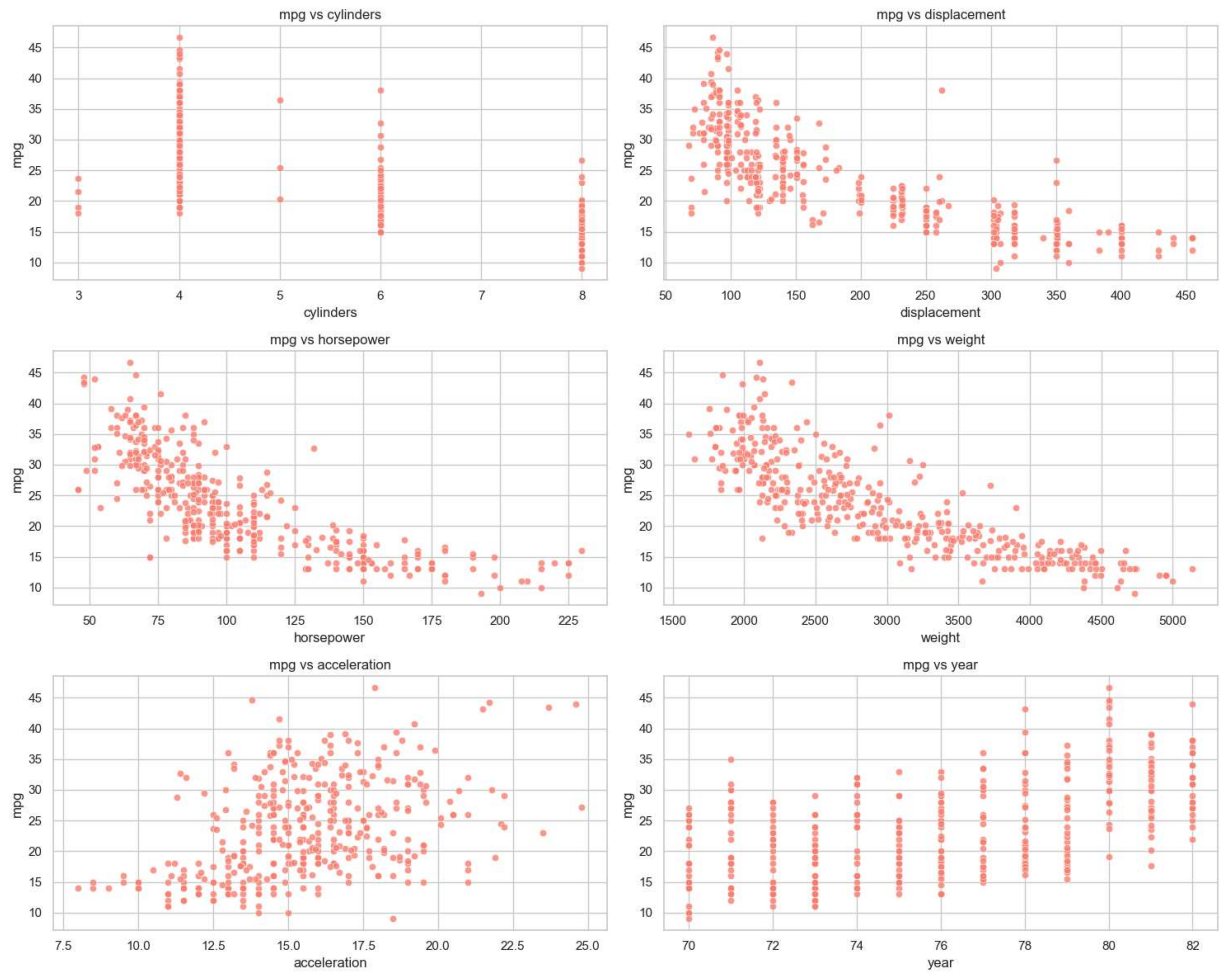
plt.tight_layout()
plt.show()
```



```
In [4]: plt.figure(figsize=(15, 12))

# Scatterplots for mpg vs other numerical features
target = 'mpg'
features = ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year']
for i, feature in enumerate(features):
    plt.subplot(3, 2, i + 1)
    sns.scatterplot(x=data[feature], y=data[target], alpha=0.8, color='salmon')
    plt.title(f'{target} vs {feature}')
    plt.xlabel(feature)
    plt.ylabel(target)

plt.tight_layout()
plt.show()
```




```
In [6]: # Correlation heatmap for numerical columns only
plt.figure(figsize=(10, 8))

# Select only numeric columns
numeric_data = data.select_dtypes(include=['float64', 'int64'])

# Calculate the correlation matrix
correlation_matrix = numeric_data.corr()

# Plot the heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Heatmap")
plt.show()
```



In []: